

Akka Streams

Franz Thoma

BOB 2018, Berlin, 2018-02-23



Streaming for Big Data Applications

Streaming Big Data

- Billions of events per day (Terabytes!)
- (Near) real-time processing
- Fault tolerance
- Bounded: Batch processing
- Unbounded: Stream processing

Why Use Akka Streams?

- Type-safe
- Compositional
- High-level
- Explicit semantics
- Integrates well (Alpakka)
- Fast (fusion & other optimizations!)

How do we use it?


- Data ingestion:
 - Real-time, stateless, CPU heavy
 - Requirements:
 - Scalable
 - Adaptable to different clients
 - Flexible deployment
 - Solution: Akka HTTP/Streams Webservices

Other Streaming Solutions


- JVM:
 - Java 8 Streams (synchronous, only trivial backpressure)
 - Java 9 Reactive Streams (rather low-level API)
 - Apache Flink (particularly for distributed systems)
- Haskell:
 - pipes
 - conduit
 - machines

Building Blocks of Akka Streams


Sources, Sinks and Flows



A Source emits (produces) items of type A



A Sink accepts (consumes) items of type A




A Flow accepts items of type A and emits items of type B.

Materialized Values

Each stream element allows to return some information on the items processed. Usually, these are information like:

- NotUsed (no information available)
- The number of elements processed
- The result (Success/Failure) of an IO action
- Items collected from the stream

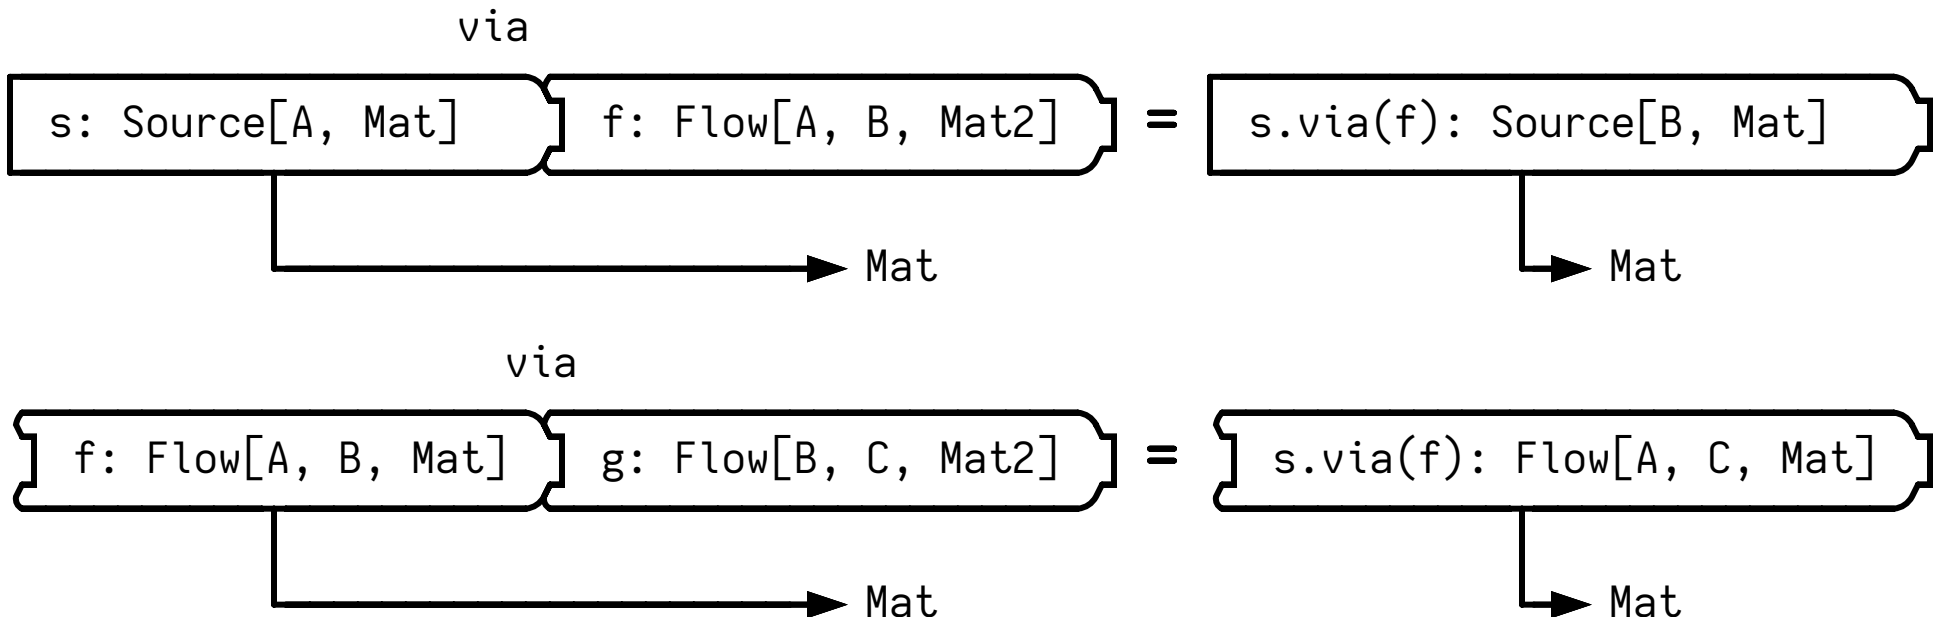
Runnable Graphs



A `RunnableGraph` is a black box that neither consumes nor produces items, but it still returns a materialized value.

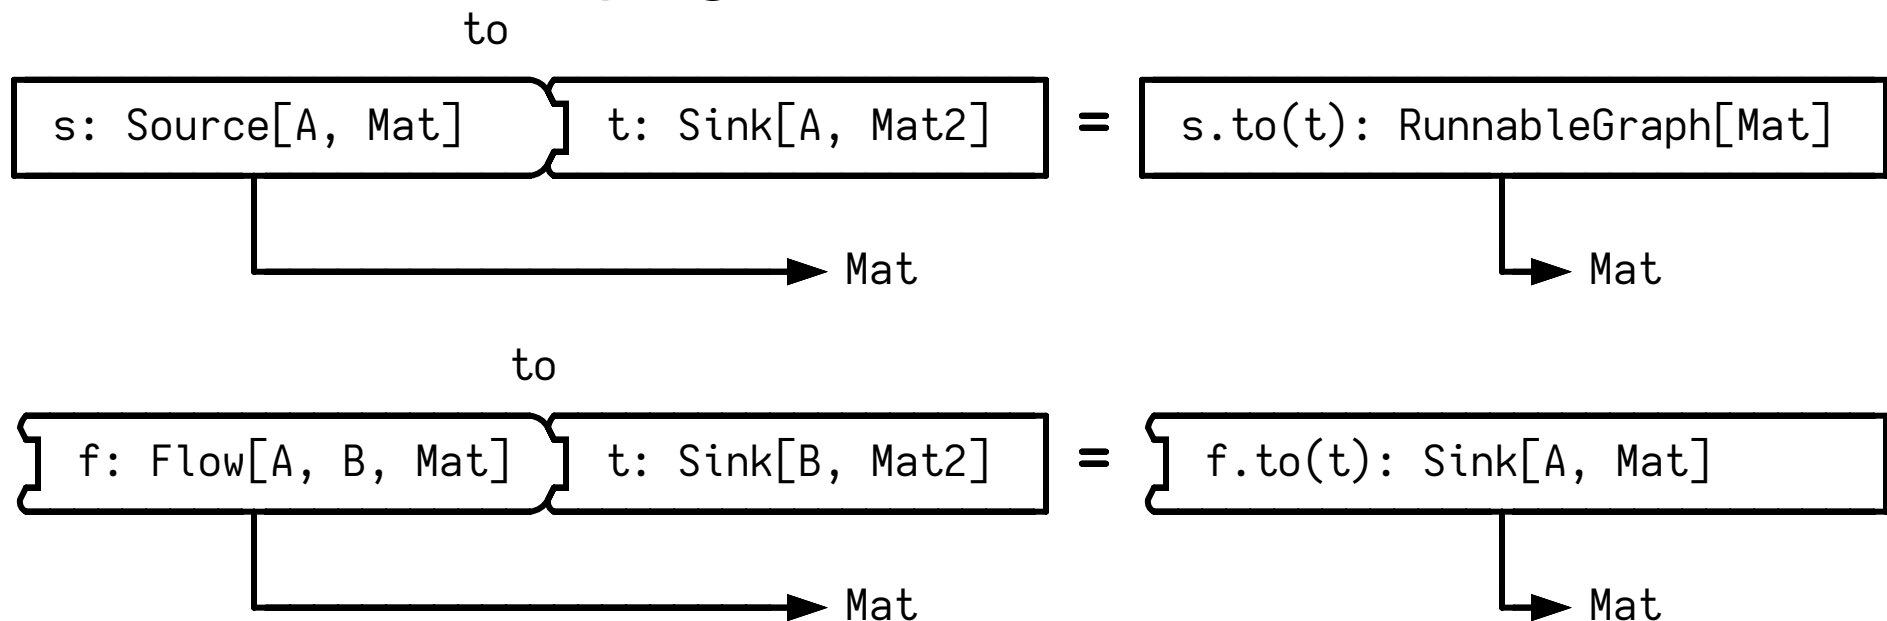
Connecting Stream Elements: *via*

via composes the the outlet of a Source or Flow with another Flow, keeping the materialized value:



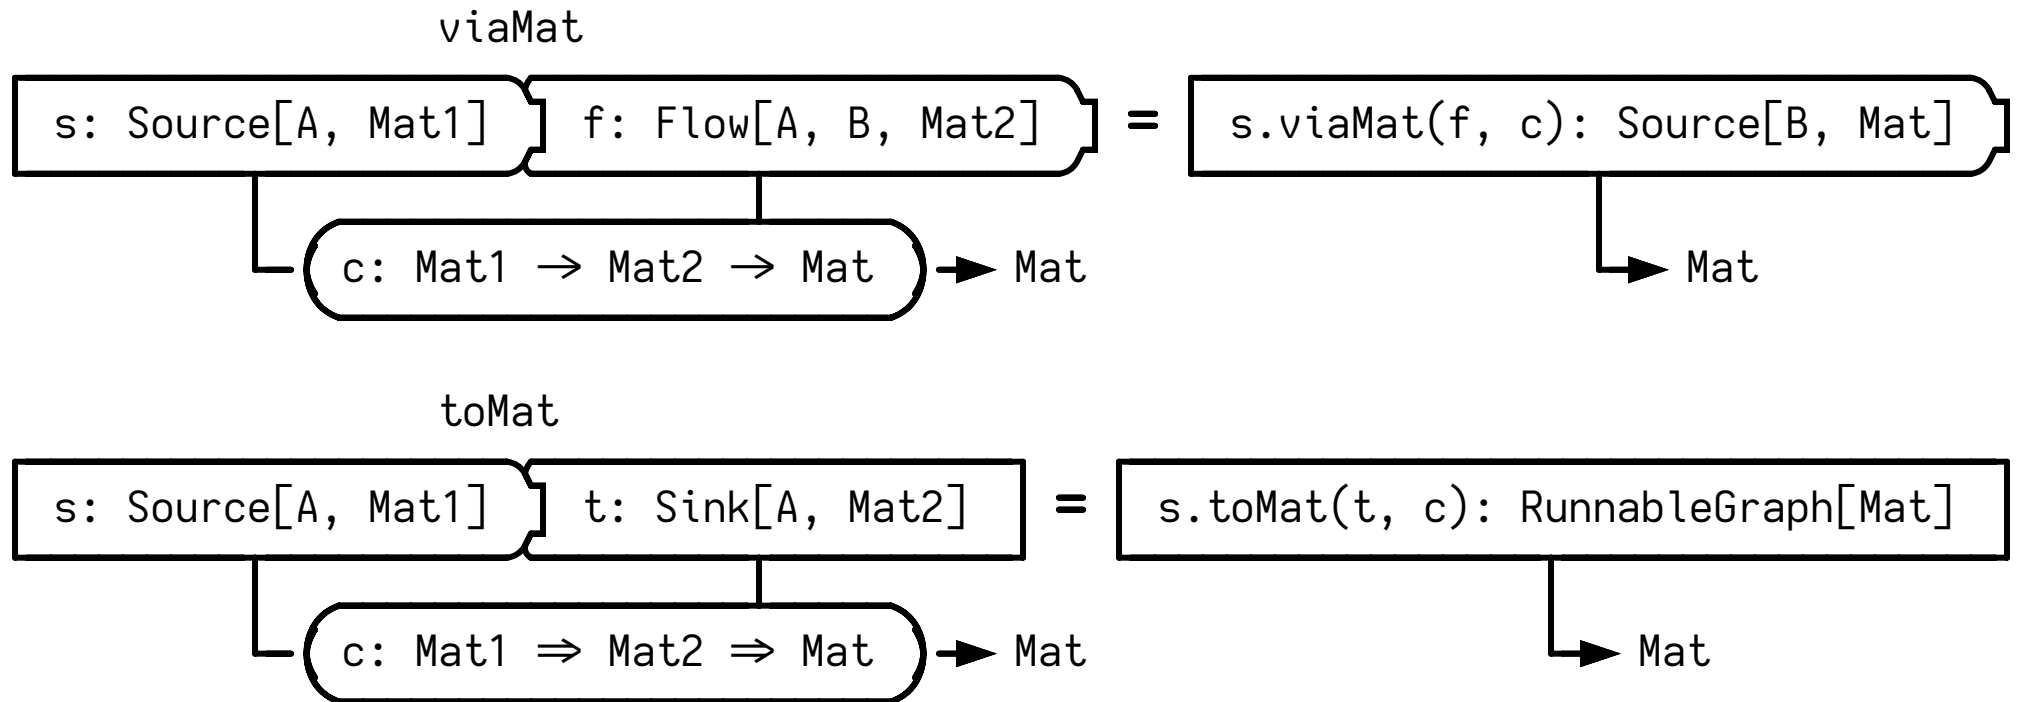
Connecting Stream Elements: to

to connects the the outlet of a Source or Flow to a Sink, keeping the materialized value:



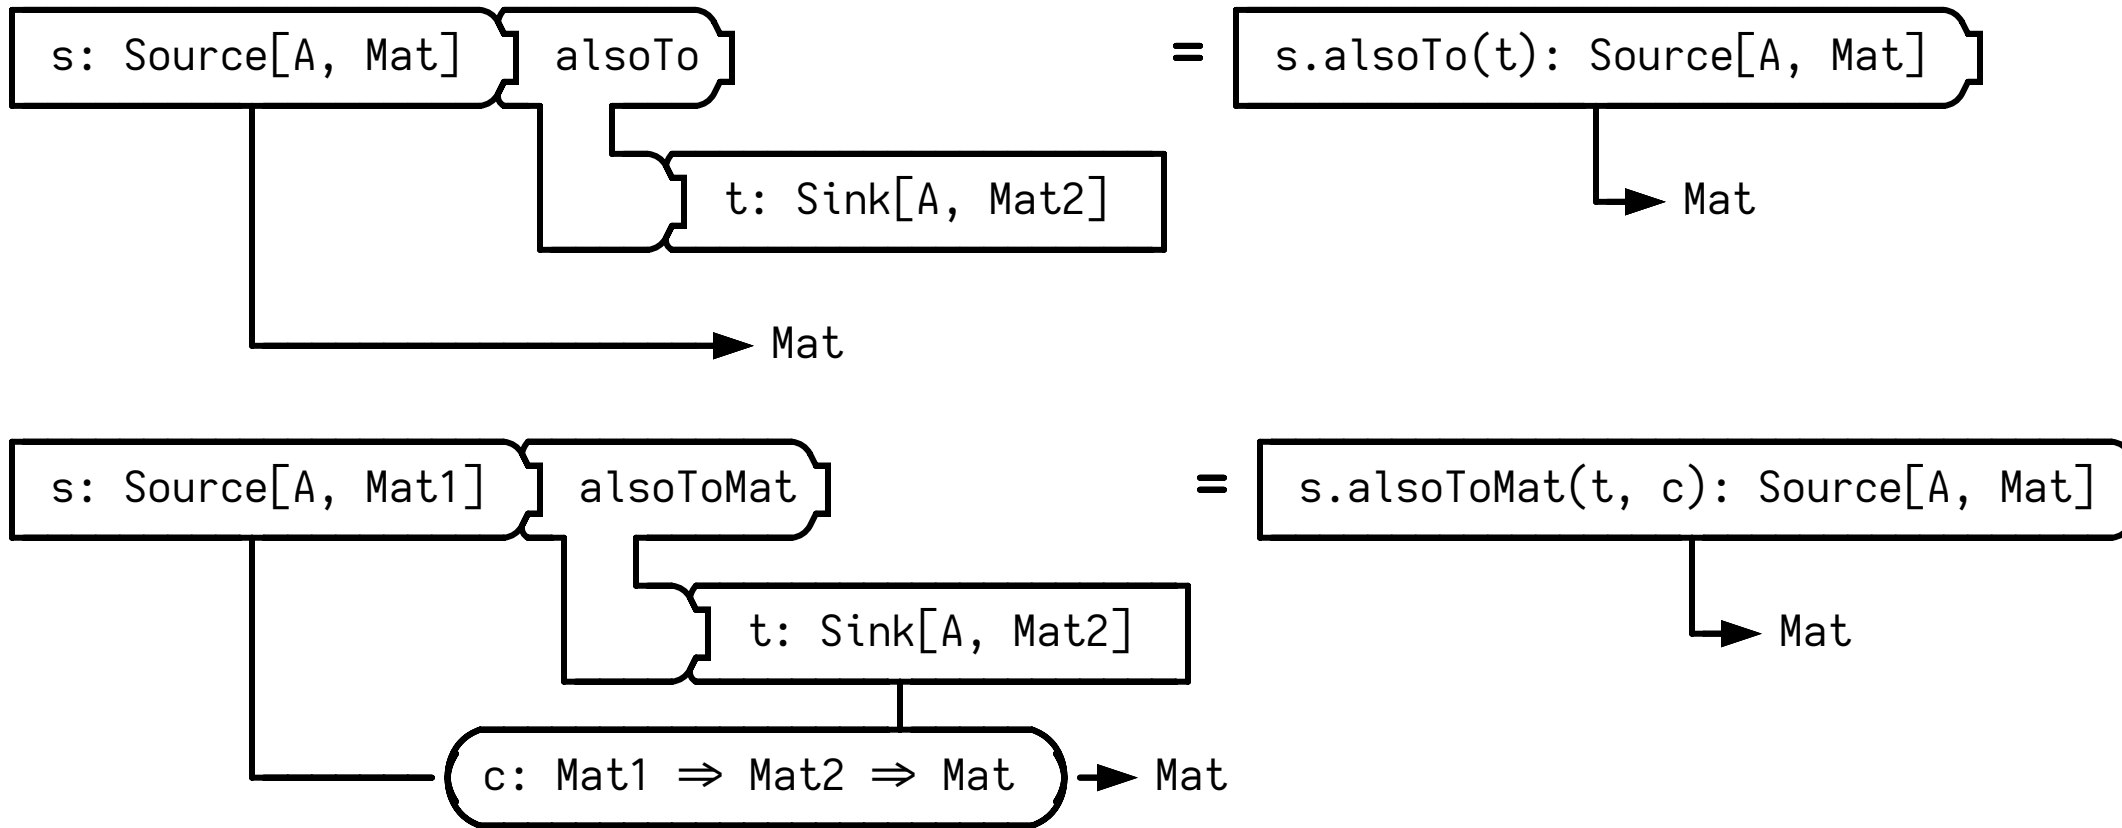
But What About the Materialized Value?

`viaMat` and `toMat` do the same, but allow to combine the materialized values:



`via(·)` is the same as `viaMat(·)` (`Keep.left`). You can also `Keep right`, `both` (returns a pair), or `none`.

Tee Pieces: alsoTo and alsoToMat



Algebraic Properties

- Source, Sink and Flow and RunnableGraph are Functors in Mat: They all have a `mapMaterializedValue` method.
- Source is a Functor in its item type: It has a method

```
map: (A ⇒ B) ⇒ Source[A, Mat] ⇒ Source[B, Mat]
```

- Sink is a contravariant Functor in its item type:

```
contramap: (A ⇒ B) ⇒ Sink[B, Mat] ⇒ Sink[A, Mat]
```

- Flow is contravariant and covariant in its both item types, respectively (aka: Profunctor): It supports both `map` and `contramap`.

Some Source, Flow and Sink examples

```
Source[T](xs: Iterable[T]): Source[T, NotUsed]

Sink.ignore: Sink[Any, Future[Done]]

Sink.foreach[T](f: T => Unit): Sink[T, Future[Done]]
// e.g. Sink.foreach(System.out.println(_))

Sink.fold[U, T](zero: U)(f: (U, T) => U): Sink[T, Future[U]]

Flow.fromFunction[A, B](f: A => B): Flow[A, B, NotUsed]

FileIO.fromPath(f: Path): Source[ByteString, Future[IOResult]]
FileIO.toPath(f: Path): Sink[ByteString, Future[IOResult]]

// Akka HTTP client/server is a Flow[HttpRequest, HttpResponse, ...]:
Http.outgoingConnection(...): Flow[HttpRequest, HttpResponse, Future[OutgoingConne
Http.bindAndHandle(handler: Flow[HttpRequest, HttpResponse, Any], ...)
```


Materialization

Sources, Sinks and Flows are just *blueprints*.

`RunnableGraph.run` builds and optimizes the actual stream.

```
implicit val system : ActorSystem = ActorSystem()
implicit val materializer : Materializer = ActorMaterializer()

val blueprint: RunnableGraph[Future[Int]] =
  Source(List(1, 2, 3)).toMat(Sink.fold(0)(_ + _))(Keep.right)

val result: Future[Int] = blueprint.run
```

Backpressure

What is Backpressure?

```
FileIO.fromPath(Paths.get("requests.txt"))           // ← fast-ish  
  .via(Framing.delimiter("\n", 1024))                // ← fast  
  .via(Fold.fromFunction(request ⇒ send(request)))   // ← slow  
  .to(Sink.foreach(response ⇒ System.out.println(response))) // ← fast
```

Default backpressuring: Only produce/consume as fast as the slowest link in the chain.

Backpressure Boundaries

What if I can't (or don't want to) control the speed of a Source?

```
incomingRequests      // ← will turn away requests if they come too fast  
  .to(slowSink)
```

```
incomingRequests      // ← won't turn away requests  
  .buffer(50, OverflowStrategy.dropNew) // ← may lose requests  
  .to(slowSink)
```

```
incomingRequests      // ← will turn away requests  
  .buffer(50, OverflowStrategy.backpressure) // if buffer is full  
  .to(slowSink)
```

Throttling

What if a Sink chokes if items come in too fast?

```
fastSource                                     // ← fast  
  .to(chokingSink)                             // ← chokes :-(
```

```
fastSource                                     // ← fast  
  .throttle(elements = 5, per = 1 second, mode = shaping) // ← slow down!  
  .to(chokingSink)                             // ← doesn't choke
```

Backpressure Boundaries (II)

```
fastSource
  .alsoTo(slowSink)      // ← slows everything down :-(
  .to(fastSink)
```

```
fastSource
  .alsoTo(Flow()
    .buffer(50, backpressure) // ← Tries to buffer
    .to(slowSink))           // ← before slowing everything down
  .to(fastSink)
```

Particularly useful if the source produces at irregular intervals.

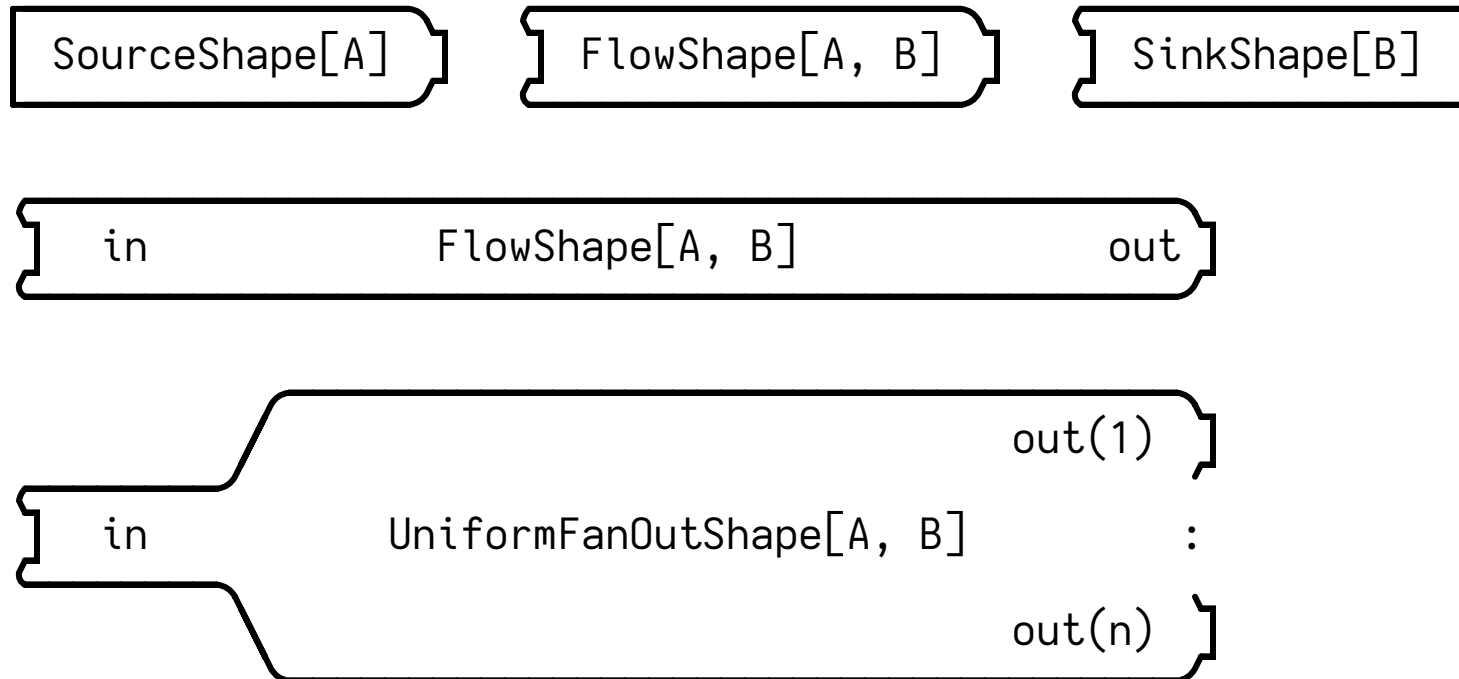
Batching

Another way to connect a fast Source to a slow Sink:

[illegible]

Graph DSL for More Complex Streams

Talk About Shapes



Shapes & Graphs

`Flow[A, B, Mat]`
is just a decorator around
`Graph[FlowShape[A, B], Mat]`

Constructing Graphs from Shapes

Balancing between workers:

```
def balanced[S, T, Mat >: Any](workers: Seq[Graph[FlowShape[S, T], Mat]]): Flow[
  Flow.fromGraph(GraphDSL.create() { implicit builder =>

    import GraphDSL.Implicits._

    val n = workers.length
    val balance: UniformFanOutShape[S, S] = builder.add(Balance[S](n))
    val merge: UniformFanInShape[T, T] = builder.add(Merge[T](n))

    for (i <- 0 until n) {
      balance.out(i) ~> workers(i).async ~> merge.in(i)
    }

    FlowShape(balance.in, merge.out)
  })
```

Also one way to speed up slow Flow elements!

Some Useful Shapes

- `SourceShape[A](Source[A, Mat])`
- `SinkShape[A](Sink[A, Mat])`
- `FlowShape[A, B](Flow[A, B, Mat])`
- `ClosedShape(RunnableGraph[Mat])`
- `FanOutShape2[A, B1, B2], FanInShape2[A1, A2, B]`
(up to 22 inlets/outlets)
- `BidiShape[In1, Out1, In2, Out2]`

Connecting to the World

Akka Streams Connectors

- Akka HTTP
- Slick (JDBC, Functional Relational Mapping)
- Apache Kafka
- Apache Camel
- AWS (S3, Kinesis, ...)
- Have a look at [Alpakka](#) for more connectors

Thank you!

Questions?

Slides on Github: [TBD](#)

[fmthoma](#) on Github

[fmthoma](#) on keybase.io

franz.thoma@tngtech.com